

125 CCs, a saving of 21.875% rather than 26.25%. Critical-word first is immune to this, since its cost is flat in the word position — which makes it the more attractive of the two when the distribution is unfavourable or unknown.

*Miss frequency (a full-AMAT caveat).* The comparison above is *per miss*, which is what this problem measures. A complete average-memory- access-time analysis would also weight each level's saving by how often that level misses, and L1 misses are far more frequent than L2 misses. That is the one factor that could push the balance back towards L1; with the numbers given here, though, L1's saving is only 3 CCs against L2's 50 CCs, so it would take an L1:L2 miss ratio above roughly 16:1 merely to draw level.